

Customer Churn Example (R)

ValidMind

2023-09-29

Analysis

Quickstart - Customer Churn Full Suite Model Documentation

This interactive notebook will guide you through documenting a model using the ValidMind Developer framework. We will use sample datasets provided by the library and train a simple classification model.

For this simple demonstration, we will use the following bank customer churn dataset from Kaggle: <https://www.kaggle.com/code/kmalit/bank-customer-churn-prediction/data>.

We will train a sample model and demonstrate the following documentation functionalities:

- Initializing the ValidMind Developer Framework
- Using a sample datasets provided by the library to train a simple classification model
- Running a test suite to quickly generate document about the data and model

Initializing the Python environment

Initializing the ValidMind Client Library

Log in to the ValidMind platform with your registered email address, and navigate to the Documentation Projects page.

Creating a new Documentation Project

(Note: if a documentation project has already been created, you can skip this section and head directly “Finding Project API key and secret”)

Clicking on “Create a new project” allows to you to register a new documentation project for our demo model.

Select “Customer Churn model” from the Model drop-down, and “Initial Validation” as Type. Finally, click on “Create Project”.

Finding the project API key and secret

In the “Client Integration” page of the newly created project, you will find the initialization code that allows the client library to associate documentation and tests with the appropriate project. The initialization code configures the following arguments:

- `api_host`: Location of the ValidMind API.
- `api_key`: Account API key.
- `api_secret`: Account Secret key.

Table 1: Numerical Variables

Name	Count	Mean	Std	Min	25%	50%	75%	90%	
RowNumber	8000	5.020520e+03	2885.7185	1	2519	5036	7512	9015	
CustomerId	8000	1.569047e+07	71902.4733	15565701	15628164	15690144	15752378	15790809	158
CreditScore	8000	6.501596e+02	96.8462	350	583	652	717	778	
Age	8000	3.894890e+01	10.4590	18	32	37	44	53	
Tenure	8000	5.033900e+00	2.8853	0	3	5	8	9	
Balance	8000	7.643410e+04	62612.2513	0	0	97264	128045	149545	16
NumOfProducts	8000	1.532500e+00	0.5805	1	1	1	2	2	
HasCrCard	8000	7.026000e-01	0.4571	0	0	1	1	1	
IsActiveMember	8000	5.199000e-01	0.4996	0	0	1	1	1	
EstimatedSalary	8000	9.979019e+04	57520.5089	12	50857	99505	149216	179486	18

- project: The project identifier. The ‘project’ argument is mandatory since it allows the library to associate all data collected with a specific account project.

The code snippet can be copied and pasted directly in the cell below to initialize the ValidMind Developer Framework when run:

```
vm_r <- vm(
  api_key="b34dfe4dcb5491212be3eefe77c85cd6",
  api_secret="40f8d2d583baa9e730a7f8872dd57e2f4657c7918c13fa259ba7ccda8a60e858",
  project="clmp6k8e800ds19mot0zu8o34",
  python_version=python_version,
  api_host="https://api.dev.vm.validmind.ai/api/v1/tracking"
)
```

Load the Demo Dataset

For the purpose of this demonstration, we will use a sample dataset provided by the ValidMind library.

```
# Read the dataset
data <- read.csv('../datasets/bank_customer_churn.csv')
```

Initialize a dataset object for ValidMind Before running the test plan, we must first initialize a ValidMind dataset object using the `init_dataset` function from the `vm` module. This function takes in arguments: `dataset` which is the dataset that we want to analyze; `target_column` which is used to identify the target variable; `class_labels` which is used to identify the labels used for classification model training.

```
vm_dataset = vm_r$init_dataset(
  dataset=data,
  target_column="Exited",
  class_labels=list("0" = "Did not exit", "1" = "Exited")
)
```

Run the Data Validation Test Plan

```
tabular_suite_results <- vm_r$run_test_suite("tabular_dataset", dataset=vm_dataset)

## HBox(children=(Label(value='Running test suite...'), IntProgress(value=0, max=24)))
raw_summary_results <- summarize_results(tabular_suite_results)
```

Table 2: Categorical Variables

Name	Count	Number of Unique Values	Top Value	Top Value Frequency	Top Value Frequency %
Surname	8000	2616	Smith	29	0.3625
Geography	8000	3	France	4010	50.1250
Gender	8000	2	Male	4396	54.9500

Table 3: Class Imbalance Results for Column Exited

Exited	Percentage of Rows (%)	Pass/Fail
0	79.80%	Pass
1	20.20%	Pass

Table 4: High Cardinality Results for Dataset

Column	Number of Distinct Values	Percentage of Distinct Values (%)	Pass/Fail
Geography	3	0.0375	Pass
Gender	2	0.0250	Pass
NumOfProducts	4	0.0500	Pass
HasCrCard	2	0.0250	Pass
IsActiveMember	2	0.0250	Pass
Exited	2	0.0250	Pass

Table 5: High Pearson Correlation Results for Dataset

Columns	Coefficient	Pass/Fail
(Balance, NumOfProducts)	-0.3044646	Fail
(Age, Exited)	0.2810129	Pass
(Exited, IsActiveMember)	-0.1515037	Pass
(Balance, Exited)	0.1174399	Pass
(Age, IsActiveMember)	0.0873114	Pass
(Exited, NumOfProducts)	-0.0523056	Pass
(Age, NumOfProducts)	-0.0306226	Pass
(CreditScore, IsActiveMember)	0.0305541	Pass
(IsActiveMember, Tenure)	-0.0293018	Pass
(Age, Balance)	0.0289801	Pass

Table 6: Missing Values Results for Dataset

Column	Number of Missing Values	Percentage of Missing Values (%)	Pass/Fail
RowNumber	0	0	Pass
CustomerId	0	0	Pass
Surname	0	0	Pass
CreditScore	0	0	Pass
Geography	0	0	Pass
Gender	0	0	Pass
Age	0	0	Pass
Tenure	0	0	Pass
Balance	0	0	Pass
NumOfProducts	0	0	Pass
HasCrCard	0	0	Pass
IsActiveMember	0	0	Pass
EstimatedSalary	0	0	Pass
Exited	0	0	Pass

Table 7: Skewness Results for Dataset

Column	Skewness	Pass/Fail
RowNumber	-0.0059207	Pass
CustomerId	0.0100323	Pass
CreditScore	-0.0619516	Pass
Age	1.0245221	Fail
Tenure	0.0076920	Pass
Balance	-0.1352769	Pass
EstimatedSalary	0.0095104	Pass

Table 8: Unique Rows Results for Dataset

Column	Number of Unique Values	Percentage of Unique Values (%)	Pass/Fail
RowNumber	8000	100.0000	Fail
CustomerId	8000	100.0000	Fail
Surname	2616	32.7000	Pass
CreditScore	452	5.6500	Pass
Geography	3	0.0375	Pass
Gender	2	0.0250	Pass
Age	69	0.8625	Pass
Tenure	11	0.1375	Pass
Balance	5088	63.6000	Pass
NumOfProducts	4	0.0500	Pass
HasCrCard	2	0.0250	Pass
IsActiveMember	2	0.0250	Pass
EstimatedSalary	8000	100.0000	Fail
Exited	2	0.0250	Pass

Table 9: Zeros Results for Dataset

Column	Number of Zero Values	Percentage of Zero Values (%)	Pass/Fail
Tenure	323	4.0375	Fail
Balance	2912	36.4000	Fail

Table 10: Model Summary Results

term	estimate	std.error	statistic	p.value
(Intercept)	-3.6279796	0.3351288	-10.8256285	0.0000000
CreditScore	-0.0008556	0.0003737	-2.2895400	0.0220480
Age	0.0717281	0.0034578	20.7438009	0.0000000
Tenure	-0.0167515	0.0124169	-1.3490929	0.1773071
Balance	0.0000017	0.0000007	2.4866260	0.0128961
NumOfProducts	-0.1089455	0.0637106	-1.7100051	0.0872649
HasCrCard	-0.0409089	0.0790667	-0.5173971	0.6048790
IsActiveMember	-1.0521795	0.0769549	-13.6726793	0.0000000
EstimatedSalary	0.0000005	0.0000006	0.7195061	0.4718291
GeographyFrance	-0.0541293	0.0944621	-0.5730267	0.5666266
GeographyGermany	0.7316662	0.1054571	6.9380484	0.0000000
GeographySpain	NA	NA	NA	NA
GenderFemale	0.4824388	0.0726963	6.6363583	0.0000000
GenderMale	NA	NA	NA	NA

Run the Model Validation Test Plan

We will need to preprocess the dataset and produce the training, test and validation splits first.

Preprocess the Raw Dataset

```
# Handle categorical variables using one-hot encoding and remove unnecessary columns
data <- data %>% select(-RowNumber, -CustomerId, -Surname)
geo_dummies <- model.matrix(~Geography - 1, data=data)
gender_dummies <- model.matrix(~Gender - 1, data=data)
data <- data %>% select(-Geography, -Gender)
data <- cbind(data, geo_dummies, gender_dummies)
```

```
# Split the dataset into training and testing sets
set.seed(123) # Setting seed for reproducibility
split <- sample.split(data$Exited, SplitRatio = 0.7)
train_data <- subset(data, split == TRUE)
test_data <- subset(data, split == FALSE)
```

```
# Run a logistic regression model
model <- glm(Exited ~ ., family = binomial(link = 'logit'), data = train_data)
```

```
model_path = save_model(model)
```

```
vm_train_ds = vm_r$init_dataset(
  dataset=train_data,
  target_column="Exited"
)
```

```
vm_test_ds = vm_r$init_dataset(
  dataset=test_data,
  target_column="Exited"
)
```

```
vm_model = vm_r$init_r_model(
  model_path=model_path,
```

Table 11: NULL

Attribute	Value
Modeling Technique	glm.fit (Family: binomial, Link function: logit)
Modeling Framework	R
Framework Version	N/A
Programming Language	R version 4.3.1 (2023-06-16)

Table 12: NULL

Dataset	Size	Proportion
Training	5600	70.00%
Test	2400	30.00%
Total	8000	100%

```

train_ds=vm_train_ds,
test_ds=vm_test_ds,
)

```

Run the Binary Classification Test Plan

```

model_validation_results = vm_r$run_test_suite("classifier_model_validation", model=vm_model, fail_fast=TRUE)

## HBox(children=(Label(value='Running test suite...'), IntProgress(value=0, max=36)))

raw_model_results <- summarize_results(model_validation_results)

```

Table 13: Classification Report

Metric	Value
Precision	0.6214099
Recall	0.2104332
F1	0.3143989
Accuracy	0.8146429
ROC AUC	0.5889938

Table 14: Classification Report

Metric	Value
Precision	0.6153846
Recall	0.1814433
F1	0.2802548
Accuracy	0.8116667
ROC AUC	0.5763613

Table 15: Population Stability Index for Training and Test Datasets

Bin	Count Initial	Percent Initial (%)	Count New	Percent New (%)	PSI
0	1845	32.9464286	825	34.3750000	0.0006064
1	1570	28.0357143	664	27.6666667	0.0000489
2	900	16.0714286	403	16.7916667	0.0003157
3	501	8.9464286	213	8.8750000	0.0000057
4	305	5.4464286	124	5.1666667	0.0001475
5	230	4.1071429	76	3.1666667	0.0024457
6	131	2.3392857	51	2.1250000	0.0002059
7	73	1.3035714	28	1.1666667	0.0001519
8	37	0.6607143	12	0.5000000	0.0004479
9	8	0.1428571	4	0.1666667	0.0000367

Table 16: Minimum Accuracy Test on Test Data

Score	Threshold	Pass/Fail
0.8116667	0.7	Pass

Table 17: Minimum F1 Score Test

Score	Threshold	Pass/Fail
0.2802548	0.5	Fail

Table 18: Minimum ROC AUC Score Test

Score	Threshold	Pass/Fail
0.5763613	0.5	Pass

Table 19: Training-Test Degradation Test

Metric	Train Score	Test Score	Degradation (%)	Pass/Fail
Accuracy	0.8146429	0.8116667	0.3653368	Pass
Precision	0.8146429	0.8116667	0.3653368	Pass
Recall	0.8146429	0.8116667	0.3653368	Pass
F1	0.8146429	0.8116667	0.3653368	Pass

Table 20: Overfit Regions Data

slice	training records	feature	training accuracy	test accuracy	gap
(349.5, 400.0]	10	CreditScore	20.000000	0.00000	20.000000
(17.926, 25.4]	321	Age	94.080997	89.03226	5.048739
(47.6, 55.0]	504	Age	60.317460	54.76190	5.555556
(77.2, 84.6]	12	Age	8.333333	0.00000	8.333333
(6.0, 7.0]	558	Tenure	86.559140	80.82707	5.732072
(25089.809, 50179.618]	35	Balance	85.714286	63.15789	22.556391
(125449.045, 150538.854]	1015	Balance	80.886700	76.55860	4.328096
(175628.663, 200718.472]	113	Balance	78.761062	63.15789	15.603167
(200718.472, 225808.281]	13	Balance	53.846154	30.76923	23.076923
(-188.338, 20003.339]	566	EstimatedSalary	82.862191	78.57143	4.290762

Table 21: Weakspots Test

slice	shape	feature	accuracy	precision	recall	f1	Dataset Type
(349.5, 400.0]	10	CreditScore	0.2000000	1.0000000	0.2000000	0.3333333	Training
(400.0, 450.0]	89	CreditScore	0.8202247	0.7142857	0.4545455	0.5555556	Training
(450.0, 500.0]	249	CreditScore	0.8152610	0.7200000	0.3157895	0.4390244	Training
(500.0, 550.0]	595	CreditScore	0.7949580	0.6666667	0.2446043	0.3578947	Training
(550.0, 600.0]	827	CreditScore	0.8210399	0.7166667	0.2471264	0.3675214	Training
(600.0, 650.0]	1020	CreditScore	0.8107843	0.5468750	0.1758794	0.2661597	Training
(650.0, 700.0]	1084	CreditScore	0.8256458	0.5409836	0.1701031	0.2588235	Training
(700.0, 750.0]	828	CreditScore	0.8200483	0.5909091	0.1656051	0.2587065	Training
(750.0, 800.0]	516	CreditScore	0.8217054	0.6750000	0.2547170	0.3698630	Training
(800.0, 850.0]	382	CreditScore	0.8036649	0.4545455	0.1369863	0.2105263	Training
(349.5, 400.0]	4	CreditScore	0.0000000	0.0000000	0.0000000	0.0000000	Test
(400.0, 450.0]	46	CreditScore	0.8043478	0.8000000	0.3333333	0.4705882	Test
(450.0, 500.0]	118	CreditScore	0.8474576	0.4285714	0.1764706	0.2500000	Test
(500.0, 550.0]	214	CreditScore	0.8224299	0.4666667	0.1891892	0.2692308	Test
(550.0, 600.0]	317	CreditScore	0.8233438	0.6875000	0.1774194	0.2820513	Test
(600.0, 650.0]	467	CreditScore	0.8094218	0.6538462	0.1752577	0.2764228	Test
(650.0, 700.0]	480	CreditScore	0.7895833	0.5312500	0.1650485	0.2518519	Test
(700.0, 750.0]	385	CreditScore	0.8000000	0.7647059	0.1511628	0.2524272	Test
(750.0, 800.0]	230	CreditScore	0.8347826	0.5789474	0.2682927	0.3666667	Test
(800.0, 850.0]	139	CreditScore	0.8417266	0.8333333	0.1923077	0.3125000	Test
(17.926, 25.4]	321	Age	0.9408100	0.0000000	0.0000000	0.0000000	Training
(25.4, 32.8]	1213	Age	0.9216818	0.0000000	0.0000000	0.0000000	Training
(32.8, 40.2]	2005	Age	0.8708229	0.6666667	0.0076923	0.0152091	Training
(40.2, 47.6]	1101	Age	0.7275204	0.6376812	0.1379310	0.2268041	Training
(47.6, 55.0]	504	Age	0.6031746	0.8301887	0.3259259	0.4680851	Training
(55.0, 62.4]	256	Age	0.7226562	0.8488372	0.5572519	0.6728111	Training
(62.4, 69.8]	111	Age	0.7567568	0.5531915	0.8125000	0.6582278	Training
(69.8, 77.2]	75	Age	0.2800000	0.0689655	1.0000000	0.1290323	Training
(77.2, 84.6]	12	Age	0.0833333	0.0833333	1.0000000	0.1538462	Training
(84.6, 92.0]	2	Age	0.0000000	0.0000000	0.0000000	0.0000000	Training
(17.926, 25.4]	155	Age	0.8903226	0.0000000	0.0000000	0.0000000	Test
(25.4, 32.8]	519	Age	0.9248555	0.0000000	0.0000000	0.0000000	Test
(32.8, 40.2]	913	Age	0.8707558	0.0000000	0.0000000	0.0000000	Test
(40.2, 47.6]	419	Age	0.6992840	0.5833333	0.1076923	0.1818182	Test
(47.6, 55.0]	210	Age	0.5476190	0.7435897	0.2543860	0.3790850	Test
(55.0, 62.4]	101	Age	0.7425743	0.7894737	0.6250000	0.6976744	Test
(62.4, 69.8]	49	Age	0.8367347	0.6250000	0.8333333	0.7142857	Test
(69.8, 77.2]	29	Age	0.3793103	0.2380952	0.7142857	0.3571429	Test
(77.2, 84.6]	3	Age	0.0000000	0.0000000	0.0000000	0.0000000	Test
(84.6, 92.0]	2	Age	0.0000000	0.0000000	0.0000000	0.0000000	Test
(-0.01, 1.0]	793	Tenure	0.8108449	0.7285714	0.2802198	0.4047619	Training
(1.0, 2.0]	606	Tenure	0.8168317	0.6000000	0.2479339	0.3508772	Training
(2.0, 3.0]	586	Tenure	0.8071672	0.5000000	0.1592920	0.2416107	Training
(3.0, 4.0]	545	Tenure	0.8036697	0.6969697	0.1916667	0.3006536	Training
(4.0, 5.0]	545	Tenure	0.8128440	0.4166667	0.1562500	0.2272727	Training
(5.0, 6.0]	547	Tenure	0.8025594	0.6451613	0.1709402	0.2702703	Training
(6.0, 7.0]	558	Tenure	0.8655914	0.7058824	0.2696629	0.3902439	Training
(7.0, 8.0]	576	Tenure	0.8246528	0.7000000	0.1858407	0.2937063	Training
(8.0, 9.0]	557	Tenure	0.8007181	0.5833333	0.1794872	0.2745098	Training
(9.0, 10.0]	287	Tenure	0.7909408	0.5555556	0.2380952	0.3333333	Training
(-0.01, 1.0]	329	Tenure	0.7933131	0.5357143	0.2142857	0.3061224	Test
(1.0, 2.0]	235	Tenure	0.8382979	0.4285714	0.0810811	0.1363636	Test
(2.0, 3.0]	243	Tenure	0.7860082	0.7142857	0.2459016	0.3658537	Test
(3.0, 4.0]	239	Tenure	0.8075314	0.6666667	0.1960784	0.3030303	Test
(4.0, 5.0]	264	Tenure	0.7878788	0.6818182	0.2343750	0.3488372	Test

Table 22: Robustness test

Perturbation Size	Dataset Type	Records	accuracy	Passed
0.0	Training	5600	81.46429	TRUE
0.0	Test	2400	81.16667	TRUE
0.1	Training	5600	81.55357	TRUE
0.1	Test	2400	81.08333	TRUE
0.2	Training	5600	81.55357	TRUE
0.2	Test	2400	80.70833	TRUE
0.3	Training	5600	81.16071	TRUE
0.3	Test	2400	80.37500	TRUE
0.4	Training	5600	80.80357	TRUE
0.4	Test	2400	80.29167	TRUE
0.5	Training	5600	80.58929	TRUE
0.5	Test	2400	80.70833	TRUE